

Lab: Greedy Best-First Search (GBFS) vs A*

Implement, analyze, and compare informed search with real maps and heuristics

Learning Goals

- Implement GBFS and A* using a priority queue.
- Design and test heuristics (admissible vs. inconsistent).
- Compare behavior: optimality, node expansions, time/space.
- Handle ties, unreachable goals, and repeated-state detection.

Problem Scenario

You're routing a drone from a start city to a goal city. You will run both GBFS (minimize $h(n)$) and A* (minimize $f(n)=g(n)+h(n)$) on two small maps:

- Map 1 (EuclidTown): coordinates are provided, use straight-line distance to the goal as $h(n)$.
- Map 2 (TrapVille): same heuristic, but edges are tweaked so GBFS is tempted down an expensive path.

Data:

Map 1: EuclidTown

Nodes & coordinates (for h to goal G):

A: (0,0) B: (2,1) C: (4,0) D: (1,3)
E: (3,4) F: (5,3) G: (6,1)

Undirected edges with costs:

A-B: 2.2 A-D: 3.2
B-C: 2.2 B-E: 3.6
C-F: 3.6
D-E: 2.4
E-F: 2.4
F-G: 2.2

Start = A, Goal = G

Reference results (for instructor checking):

Optimal path (by Dijkstra): $A \rightarrow B \rightarrow C \rightarrow F \rightarrow G$

Optimal cost: 10.2

With straight-line h to G , A^* is optimal; GBFS also finds the same path here.

Map 2: TrapVille

Nodes & coordinates:

$S:(0,0)$ $X:(2,2)$ $Y:(2,-2)$ $Z:(4,0)$ $G:(6,0)$

Undirected edges with costs:

$S-X:2.9$ $S-Y:2.9$

$X-Z:8.0$ $Y-Z:2.6$

$Z-G:2.0$

Start = S , Goal = G

Reference results:

True optimal: $S \rightarrow Y \rightarrow Z \rightarrow G$, cost 7.5

GBFS (typical ties) is lured into: $S \rightarrow X \rightarrow Z \rightarrow G$, cost 12.9 (suboptimal)

A^* returns the optimal 7.5

Required Functions:

- `priority_queue_push(pq, priority, state, g_cost, parent)`
- `priority_queue_pop(pq) -> (priority, state, g_cost, parent)`
- `gbfs(start, goal, graph, heuristic) -> (path, cost, expansions, visited_order)` — priority = $h(n)$
- `astar(start, goal, graph, heuristic) -> (path, cost, expansions, visited_order)` — priority = $g(n)+h(n)$ and maintain best- g
- `reconstruct_path(parents, goal) -> list_of_nodes`
- `euclidean_heuristic(node, goal, coords) -> float`
- `is_consistent(graph, heuristic) -> (bool, violating_edges[])` — check $h(u) \leq \text{cost}(u,v)+h(v)$ for all edges
- `is_admissible(all_nodes, goal, graph, heuristic) -> (bool, counterexamples[])` — compare to true shortest path

Note: If heaps are confusing, you may use a sorted list/array (slower but acceptable for small graphs). Document your choice.

What to Submit

Part A Correctness (EuclidTown)

- Run A^* on Map 1: report path, total cost, nodes expanded, visited order.

- Run GBFS on Map 1: same reports.
- Verify both return $A \rightarrow B \rightarrow C \rightarrow F \rightarrow G$ with cost ≈ 10.2 (minor float error OK).
- Show `is_consistent == True` and `is_admissible == True` for the heuristic.

Part B GBFS Trap (TrapVille)

- Run A* on Map 2 expect $S \rightarrow Y \rightarrow Z \rightarrow G$, cost ≈ 7.5 .
- Run GBFS on Map 2 likely $S \rightarrow X \rightarrow Z \rightarrow G$, cost ≈ 12.9 (explain tie-breaking).
- Compare expansions and explain why GBFS fails (it ignores $g(n)$).

Part C Tie- Breaking & Variants

- Change tie-break rule (lexical vs FIFO vs LIFO on equal priority) and re-run GBFS on Map 2. Did results change? Why?
- (Optional bonus) Add a slightly noisy heuristic $h'(n) = h(n) * k$ where $k \in [0.95, 1.05]$ and show sensitivity for GBFS vs A*.

Marking Rubric (30 Marks)

Implementation (12)	GBFS priority logic (3); A* with best-g & open/closed (6); Path reconstruction & I/O (3)
Heuristic Checks (6)	<code>is_consistent</code> (3); <code>is_admissible</code> (3)
Experiments & Results (8)	Map 1 correct & explained (3); Map 2 trap & analysis (3); Tie-breaking (2)
Clarity & Code Quality (4)	Comments, modularity, clean logs
Optional Bonus (+3)	Noisy-heuristic sensitivity with concise charts/tables